



AMALTHEA-based GPU Response Time Analysis for NVidia's Jetson TX2

Daniel Paredes



Agenda:


1. Motivation
2. Nvidia GPU Model
3. AMALTHEA Model
4. Response Time Analysis for Jetson TX2's GPU
5. Results
- 6. Conclusions**

Motivation

Motivation

WATERS Industrial Challenge:

- Heterogeneous systems
- Centralized end-to-end architectures.
- Chosen platform: NVIDIA Jetson™ TX2



NVidia GPU Model



Nvidia GPU Model

Software Model:


- Threads
- Blocks
- Grids

Hardware Model:

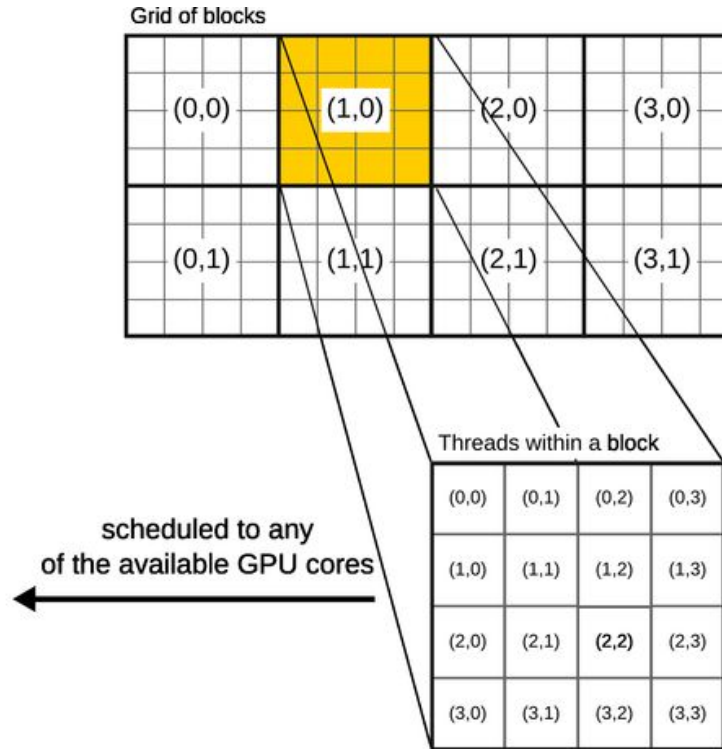
- Streaming Multiprocessors
- Memory

Nvidia GPU Model

Streaming Multiprocessors



(a)

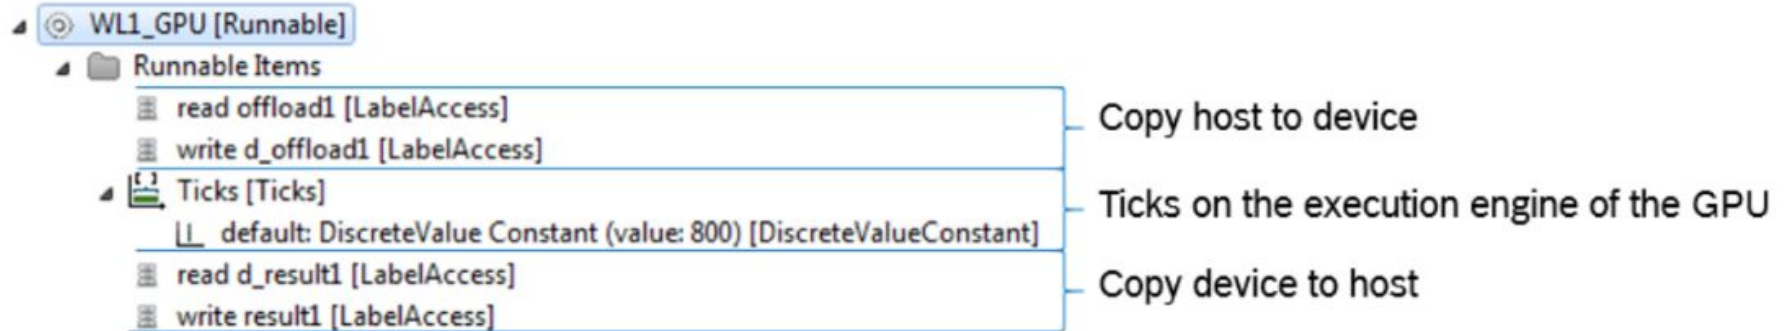


(b)

[1] Computing Efficiently Spectral-Spatial Classification of Hyperspectral Images on Commodity GPUs - Pablo Quesada-Barriuso

AMALTHEA Model

AMALTHEA Model



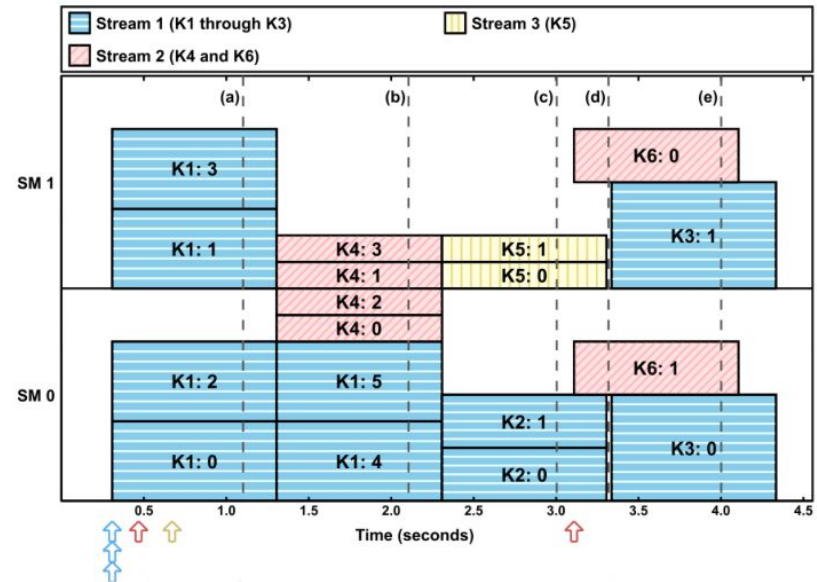
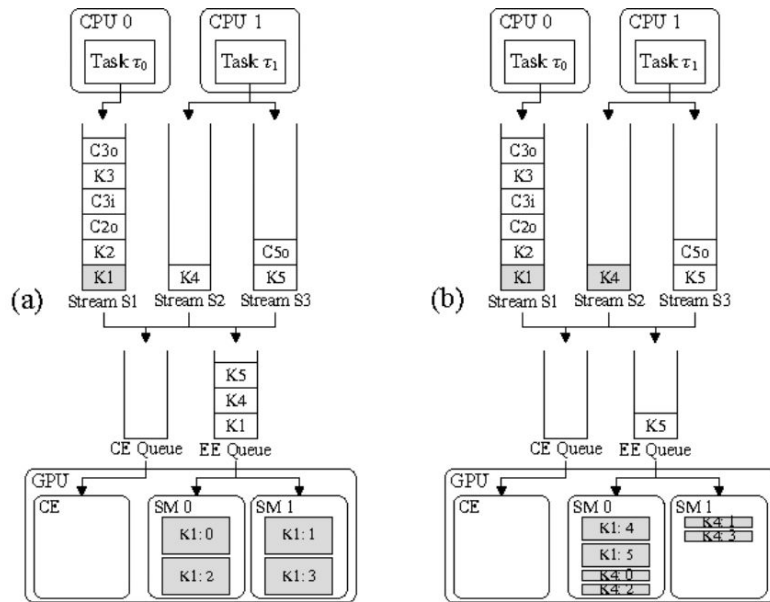
Response Time Analysis



Jetson TX2's GPU


- 2 Streaming Multiprocessors
- FIFO Queues
- Max. threads per Block: 1024
- Max. threads per SM: 2048

GPU Scheduler




[2] GPU Scheduling on the NVIDIA TX2: Hidden Details Revealed - Tanya Amert et al.

Response Time Calculation






Response Time Calculation




(a)

Response Time Calculation




(a)




(b)


Response Time Calculation



(a)



(b)



(c)



Algorithm

τ : set of kernels

f_i : completion time of kernel i

C_i : thread execution time


Input : τ

Output: f_1, \dots, f_n

Initialization: $t_a = 0, g_f = g_{max}, i = 1$

```
while  $i \leq n$  do
    if  $g_f \geq g_i$  then
         $f_i = t_a + C_i$ ;
        Update  $g_f$  and  $t_a$ ;
         $i++$  ; // Next kernel
    else
         $g_i = g_i - g_f$ ;
        Update  $g_f$  and  $t_a$ ;
    end
end
```

Algorithm



Input : τ

Output: f_1, \dots, f_n

Initialization: $t_a = 0$, $g_f = g_{max}$, $i = 1$

while $i \leq n$ do

 if $g_f \geq g_i$ then

$f_i = t_a + C_i$;

 Update g_f and t_a ;

$i++$; // Next kernel

 else


$g_i = g_i - g_f$;

 Update g_f and t_a ;

 end

end

Algorithm



Input : τ

Output: f_1, \dots, f_n


Initialization: $t_a = 0$, $g_f = g_{max}$, $i = 1$, $h = \{\}$

```

while  $i \leq n$  do
    if  $g_f \geq g_i$  then
         $f_i = t_a + C_i$ ;
         $h = \{h; (f_i, g_i)\}$ ;
         $t_a = t_a$ ;
         $g_f = g_f - g_i$ ;
         $i++$  ; // Next kernel
    else
         $g_i = g_i - g_f$ ;
         $h = \{h; (t_a + C_i, g_f)\}$ ;
         $[t_a, \text{index}] = \min(h[:, 1])$ ;
         $g_f = h[\text{index}, 2]$ ;
        Update  $h$ ;
    end
end

```

Algorithm



Input : τ

Output: f_1, \dots, f_n

Initialization: $t_a = 0$, $g_f = g_{max}$, $i = 1$, $h = \{\}$

while $i \leq n$ do

if $g_f \geq g_i$ then

$f_i = t_a + C_i$;

$h = \{h; (f_i, g_i)\}$;

$t_a = t_a$;

$g_f = g_f - g_i$;

$i++$; // Next kernel

else

$g_i = g_i - g_f$;

$h = \{h; (t_a + C_i, g_f)\}$;

$[t_a, \text{index}] = \min(h[:, 1])$;

$g_f = h[\text{index}, 2]$;

Update h ;

end

end

AMALTHEA implementation

```
while  $i \leq n$  do
  if  $g_f \geq g_i$  then
     $f_i = t_a + C_i$ ;
     $h = \{h; (f_i, g_i)\}$ ;
     $t_a = t_a$ ;
     $g_f = g_f - g_i$ ;
     $i++$  ; // Next kernel
  else
     $g_i = g_i - g_f$ ;
     $h = \{h; (t_a + C_i, g_f)\}$ ;
     $[t_a, \text{index}] = \min(h[:, 1])$ ;
     $g_f = h[\text{index}, 2]$ ;
    Update  $h$ ;
  end
end
```

```
// Main loop
while ( current_kernel < rList.size() ) {
  if (  $g_f \geq g_i.get(\text{current\_kernel})$  ) {
     $f.add(\text{current\_kernel}, t_a + c_i.get(\text{current\_kernel}) )$  ;

     $h = \text{updateH}(h, f.get(\text{current\_kernel}), g_i.get(\text{current\_kernel}) )$  ;

     $g_f = g_f - g_i.get(\text{current\_kernel})$ ;
     $\text{current\_kernel}++$ ;
  }
  else {
     $g_i.set(\text{current\_kernel}, g_i.get(\text{current\_kernel}) - g_f)$ ;

     $h = \text{updateH}(h, t_a + c_i.get(\text{current\_kernel}), g_f)$ ;
     $\text{minimumRegisteredTicks} = \text{findIndexOffMinValue}(h)$ ;

     $g_f = h.get(\text{minimumRegisteredTicks})$ ;
     $t_a = \text{minimumRegisteredTicks}$ ;

     $h.remove(\text{minimumRegisteredTicks})$ ;
  }
}
```

AMALTHEA implementation

```
while  $i \leq n$  do
  if  $g_f \geq g_i$  then
     $f_i = t_a + C_i$ ;
     $h = \{h; (f_i, g_i)\}$ ;
     $t_a = t_a$ ;
     $g_f = g_f - g_i$ ;
     $i++$  ; // Next kernel
  else
     $g_i = g_i - g_f$ ;
     $h = \{h; (t_a + C_i, g_f)\}$ ;
     $[t_a, \text{index}] = \min(h[:, 1])$ ;
     $g_f = h[\text{index}, 2]$ ;
    Update  $h$ ;
  end
end
```

```
// Main loop
while ( current_kernel < rList.size() ) {
  if (  $g_f \geq g_i.get(\text{current\_kernel})$  ) {
     $f.add(\text{current\_kernel}, t_a + c_i.get(\text{current\_kernel}) )$  ;

     $h = \text{updateH}(h, f.get(\text{current\_kernel}), g_i.get(\text{current\_kernel}) )$  ;

     $g_f = g_f - g_i.get(\text{current\_kernel})$ ;
     $\text{current\_kernel}++$ ;
  }
  else {
     $g_i.set(\text{current\_kernel}, g_i.get(\text{current\_kernel}) - g_f)$ ;

     $h = \text{updateH}(h, t_a + c_i.get(\text{current\_kernel}), g_f)$ ;
     $\text{minimumRegisteredTicks} = \text{findIndexOffMinValue}(h)$ ;

     $g_f = h.get(\text{minimumRegisteredTicks})$ ;
     $t_a = \text{minimumRegisteredTicks}$ ;

     $h.remove(\text{minimumRegisteredTicks})$ ;
  }
}
```

AMALTHEA implementation

```
while  $i \leq n$  do
  if  $g_f \geq g_i$  then
     $f_i = t_a + C_i$ ;
     $h = \{h; (f_i, g_i)\}$ ;
     $t_a = t_a$ ;
     $g_f = g_f - g_i$ ;
     $i++$  ; // Next kernel
  else
     $g_i = g_i - g_f$ ;
     $h = \{h; (t_a + C_i, g_f)\}$ ;
     $[t_a, \text{index}] = \min(h[:, 1])$ ;
     $g_f = h[\text{index}, 2]$ ;
    Update  $h$ ;
  end
end
```

```
// Main loop
while ( current_kernel < rList.size() ) {
  if (  $g_f \geq g_i.get(\text{current\_kernel})$  ) {
     $f.add(\text{current\_kernel}, t_a + c_i.get(\text{current\_kernel}) )$  ;

     $h = \text{updateH}(h, f.get(\text{current\_kernel}), g_i.get(\text{current\_kernel}) )$  ;

     $g_f = g_f - g_i.get(\text{current\_kernel})$ ;
     $\text{current\_kernel}++$ ;
  }
  else {
     $g_i.set(\text{current\_kernel}, g_i.get(\text{current\_kernel}) - g_f)$ ;

     $h = \text{updateH}(h, t_a + c_i.get(\text{current\_kernel}), g_f)$ ;
     $\text{minimumRegisteredTicks} = \text{findIndexOffMinValue}(h)$ ;

     $g_f = h.get(\text{minimumRegisteredTicks})$ ;
     $t_a = \text{minimumRegisteredTicks}$ ;

     $h.remove(\text{minimumRegisteredTicks})$ ;
  }
}
```

AMALTHEA implementation

```
while  $i \leq n$  do
  if  $g_f \geq g_i$  then
     $f_i = t_a + C_i$ ;
     $h = \{h; (f_i, g_i)\}$ ;
     $t_a = t_a$ ;
     $g_f = g_f - g_i$ ;
     $i++$  ; // Next kernel
  else
     $g_i = g_i - g_f$ ;
     $h = \{h; (t_a + C_i, g_f)\}$ ;
     $[t_a, \text{index}] = \min(h[:, 1])$ ;
     $g_f = h[\text{index}, 2]$ ;
    Update  $h$ ;
  end
end
```

```
// Main loop
while ( current_kernel < rList.size() ) {
  if (  $g_f \geq g_i.get(\text{current\_kernel})$  ) {
     $f.add(\text{current\_kernel}, t_a + c_i.get(\text{current\_kernel}) )$  ;

     $h = \text{updateH}(h, f.get(\text{current\_kernel}), g_i.get(\text{current\_kernel}) )$  ;

     $g_f = g_f - g_i.get(\text{current\_kernel})$ ;
     $\text{current\_kernel}++$ ;
  }
  else {
     $g_i.set(\text{current\_kernel}, g_i.get(\text{current\_kernel}) - g_f)$ ;

     $h = \text{updateH}(h, t_a + c_i.get(\text{current\_kernel}), g_f)$ ;
     $\text{minimumRegisteredTicks} = \text{findIndexOffMinValue}(h)$ ;

     $g_f = h.get(\text{minimumRegisteredTicks})$ ;
     $t_a = \text{minimumRegisteredTicks}$ ;

     $h.remove(\text{minimumRegisteredTicks})$ ;
  }
}
```


AMALTHEA implementation

```
while  $i \leq n$  do
  if  $g_f \geq g_i$  then
     $f_i = t_a + C_i$ ;
     $h = \{h; (f_i, g_i)\}$ ;
     $t_a = t_a$ ;
     $g_f = g_f - g_i$ ;
     $i++$  ; // Next kernel
  else
     $g_i = g_i - g_f$ ;
     $h = \{h; (t_a + C_i, g_f)\}$ ;
     $[t_a, \text{index}] = \min(h[:, 1])$ ;
     $g_f = h[\text{index}, 2]$ ;
    Update  $h$ ;
  end
end
```

```
// Main loop
while ( current_kernel < rList.size() ) {
  if (  $g_f \geq g_i.get(\text{current\_kernel})$  ) {
     $f.add(\text{current\_kernel}, t_a + c_i.get(\text{current\_kernel}) )$  ;

     $h = \text{updateH}(h, f.get(\text{current\_kernel}), g_i.get(\text{current\_kernel}) )$  ;

     $g_f = g_f - g_i.get(\text{current\_kernel})$ ;
     $\text{current\_kernel}++$ ;
  }
  else {
     $g_i.set(\text{current\_kernel}, g_i.get(\text{current\_kernel}) - g_f)$ ;

     $h = \text{updateH}(h, t_a + c_i.get(\text{current\_kernel}), g_f)$ ;
     $\text{minimumRegisteredTicks} = \text{findIndexOffMinValue}(h)$ ;

     $g_f = h.get(\text{minimumRegisteredTicks})$ ;
     $t_a = \text{minimumRegisteredTicks}$ ;

     $h.remove(\text{minimumRegisteredTicks})$ ;
  }
}
```


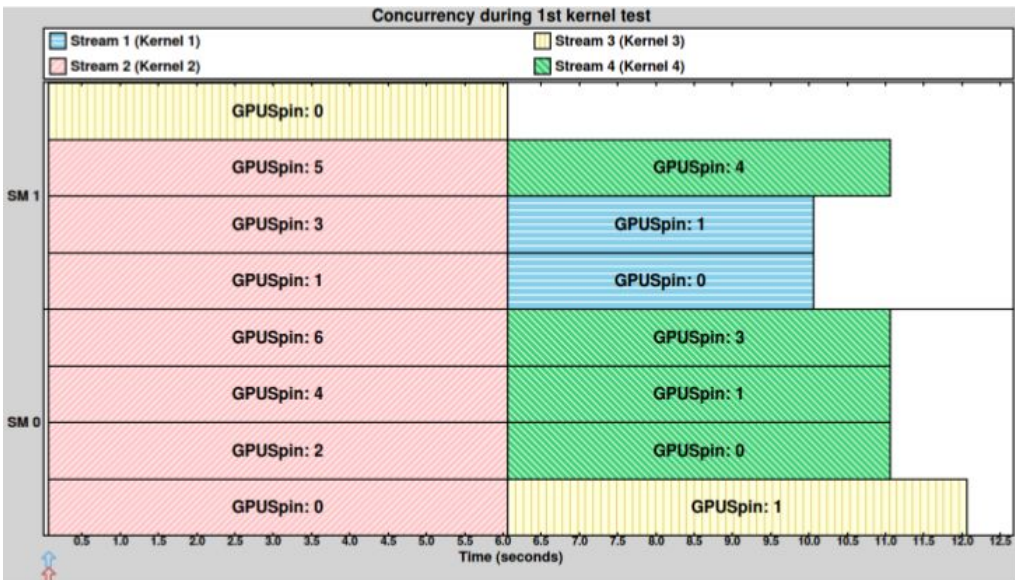
Results





Set-up

- Dummy kernels on Jetson TX2
- AMALTHEA implementation
 - Timing measurement
 - Block allocation diagrams

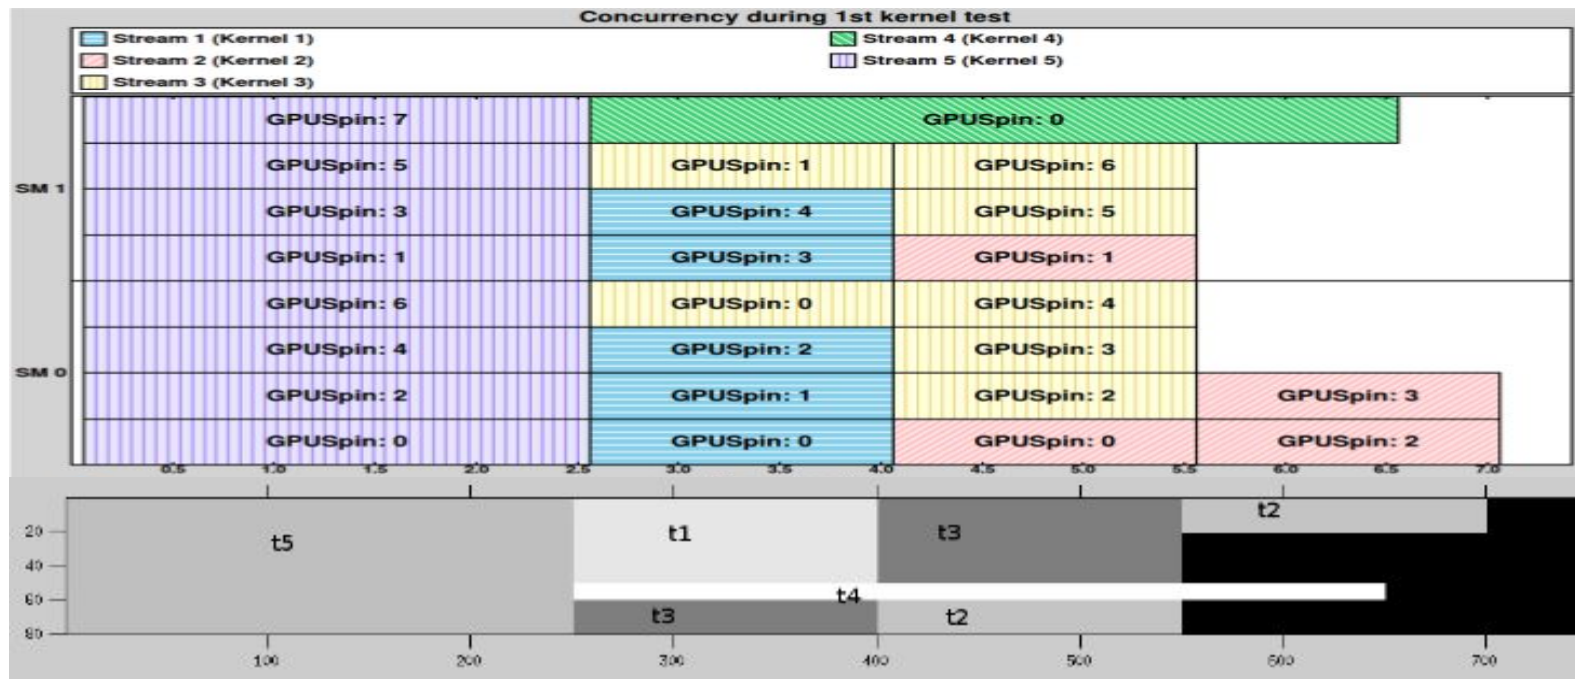
Results



Results



Results



Conclusions and Future work



Conclusions

- Simple and easy to implement algorithm
- Accuracy in calculation completion times
- Integration with AMALTHEA models
- Simulation of complex use cases on AMALTHEA models



Future Work

- Automatic CUDA code generation
- Consideration of memory transactions
- Not fully understand of Jetson TX2's GPU behaviour



References:

1. Quesada-Barriuso, Pablo & Argüello, Francisco & B. Heras, Dora. (2013). Computing Efficiently Spectral-Spatial Classification of Hyperspectral Images on Commodity GPUs. 10.1007/978-3-319-01649-8_2.
2. T. Amert, N. Otterness, M. Yang, J. H. Anderson, and F. D. Smith, "GPU Scheduling on the Nvidia TX2: Hidden Details Revealed," in 2017 IEEE Real-Time Systems Symposium (RTSS), 2017, pp. 104–115.



AMALTHEA-based GPU Response Time Analysis for NVidia's Jetson TX2

Daniel Paredes